Assessing Toxicity in Wikipedia Comments

Jonathan Innis, Gabriel Britain

Contents

Abstract

As the number of users on the internet increases, so do the number of obscene, hateful, or otherwise-toxic comments, and the need to moderate those comments. We present a text classifier model to determine whether a model is toxic or not, and if toxic, in what way. The topic was originally presented as a $35,000 competition on Kaggle, and uses data as provided by Wikipedia talk page edits [7]. A number of models were developed through trial-and-error to classify comments. Specifically developed models were Naive Bayes, support-vector machines, random forest classifiers, and a GloVe embedding-based recurrent neural network.

1   Dataset Source *Gabriel Britain*

The dataset used can be found on Kaggle under the "Toxic Comment Classification Challenge." Another participant in the challenge performed data augmentation by running the original training dataset repeatedly through Google Translate, and then including the results as training data. This augmented data, however, could lead to significant overfitting, as discussed later.

2   Frameworks *Gabriel Britain*

For data analyis and graphing, pandas[5] and matplotlib[4] were used. For the process of actual machine learning model building, rather than spending a significant amount of time implementing different models from scratch, and to focus efforts on feature engineering, architecture design, and hyperparameter optimization, scikit-learn[2] and Keras[3] were used.

2.1   Pandas

Pandas is an open-source library that provides high-performance and easy-to-use data structures for processing and analyzing CSV data. Pandas was used for data inspection, analyzing class imbalances, and reporting results of developed algorithms.

2.2   Matplotlib

Matplotlib is a Python 2D plotting library which produces high-quality figures easily. Matplotlib was used occasionally for plotting the distribution of words, characters, phrases, and class imbalances.

2.3   Scikit-Learn

Scikit-Learn (sklearn) is a simple and efficient machine learning framework that implements many machine learning algorithms and packages them in simple,

easy-to-use Python classes. Hyperparameters for these machine learning algorithms are available as Python class attributes, which makes hyperparameter tuning a very simple task. In addition, sklearn implements many functions that simplify the model-building workflow, such as cross-validation, training-validation splits, tokenization, and metric reporting, and model evaluation.

2.4 Keras

Keras is a high-level neural networks API, written in Python and capable of interfacing with TensorFlow[1]. Keras was used to develop a recurrent neural network model that used pre-trained GloVe word embeddings[6], to determine if pre-trained word embeddings could provide stronger and more robust models than statistics-based models.

3 Evaluation Metrics *Jonathan Innis*

The task of identifying toxic comments can be modeled as a multilabel classification problem. As a result, we set out to find evaluation metrics that were appropriate for the task. Initially, Hamming loss and Jaccard similarity were used as our evaluation metrics, but were found to be too forgiving. Thus, we settled on using F1 as our evaluation metric for each class. To evaluate the overall success of the model, we used two aggregate metrics: weighted F1 (which is the weighted average of F1 scores for each class), and macro F1.

As we continued to improve and evaluate our models, we became curious as to the performance of our models as compared to those entered in the original competition. We found that the original Kaggle competition was actually using Area under the Receiver Operating Curve (AUROC) as their evaluation metric for submission quality. We didn't use this metric for evaluating our model, but rather to compare the performance of our model with other submitted models.

4 Process

4.1 Data Analysis *Jonathan Innis, Gabriel Britain*

Analysis began upon downloading the training dataset. The overall class distribution can be seen in Figure 1 on page 3. The "non-toxic" label was artificially added to indicate the number of examples that are labeled as not belonging to any of the toxicity classes. Clearly, the number of "non-toxic" examples dominates the number of toxic examples belonging to any category. Within the toxic labels, however, there is *also* a class imbalance, as can be seen in Figure 2 on page 4.

The "severe_toxic", "threat", and "identity_hate" classes are significantly under-represented in the toxic classes. Further data analysis can be found in the "Data Introspection" file under the `src` folder.

Figure 1: Class Distribution

4.2 Establishing a Baseline *Gabriel Britain*

In order to evaluate the models developed and verify that they were having an actual impact on classification quality, a baseline was created using the distribution of toxicity classes found. The result of the baseline on the testing dataset can be found in Figure 1 on page 3.

|	Precision	Recall	F1-Score	Support
Toxic	0.10	0.06	0.16	5038
Severely Toxic	0.01	0.06	0.02	500
Obscene	0.05	0.23	0.08	2810
Threat	0.00	0.02	0.08	152
Insult	0.05	0.23	0.01	2591
Identity Hate	0.01	0.04	0.01	449
Micro Average	0.07	0.30	0.11	11540
Macro Average	0.04	0.17	0.06	11540
Weighted Average	*0.07*	*0.30*	*0.11*	*11540*

Table 1: Random Baseline Results

Figure 2: Toxic Class Distribution

5 Model Development

5.1 Feature Engineering *Gabriel Britain*

Initially, the only feature used in training each model was the tokenized document corpus as a bag-of-words (BOW), represented through TF-IDF vectors. However, we soon realized that there were some words that automatically marked a comment as being of a certain class. Most significantly, the presence of typical obscenities very often leads to a comment being classified as "obscene". Once we included a feature vector comprised of binary values, indicating whether an obscenity was present in the comment or not, the classification accuracy of the "obscene" category jumped significantly. The same approach was taken to identifying "identity hate" comments by including a feature vector of slurs. However, the underrepresentation of the "identity hate" class proved to be more difficult than that of the "obscene" class, and did not lead to similar results.

5.2 Naive Bayes Classifier *Jonathan Innis*

The first model built was a Naive Bayes classifier. A bag of words model is typically a strong baseline for most text classification problems. In this par-

ticular task, the presence/absence of certain words in a comment can classify a comment as being threatening, toxic, obscene, etc. Thus, if a certain word appears in a comment, it is most likely that that comment will fall into that particular category. When the results of the Naive Bayes implementation were inspected without feature optimization, the F1-scores of each class had greatly increased from the baseline. As a result, the macro and weighted averages of the F1 scores had *also* increased significantly. After performing an exhaustive hyperparameter search through sklearn's `GridSearchCV`, the overall recall of the system increased, while the overall precision decreased slightly. The results of the hyperparameter tuning can be seen in Figure 2 on page 5.

	Precision	Recall	F1-Score	Support
Toxic	0.10	0.06	0.16	5038
Severely Toxic	0.01	0.06	0.02	500
Obscene	0.05	0.23	0.08	2810
Threat	0.00	0.02	0.08	152
Insult	0.05	0.23	0.01	2591
Identity Hate	0.01	0.04	0.01	449
Micro Average	0.07	0.30	0.11	11540
Macro Average	0.04	0.17	0.06	11540
Weighted Average	*0.07*	*0.30*	*0.11*	*11540*

Table 2: Tuned Naive Bayes on Test Set

Despite the macro-averaged F1 score not changing, the weighted-average F1 score increased slightly, improving the classifier.

5.3   Support Vector Machines (SVMs) *Jonathan Innis*

The second model created was an SVM model, to see if there would be improvements over the Naive Bayes classifier. The same feature set was used as in the Naive Bayes classifier, and the same exhaustive hyperparameter tuning was performed. However, even with extensive hyperparameter tuning, the support vector machines performed significantly worse than the Naive Bayes classifier on both training and testing sets, as seen in Figure 3 on page 6.

5.4   Random Forest Classifier *Gabriel Britain*

The next model constructed was a random forest classifier. Random forest classifiers can accommodate imbalanced datasets by weighting majority classifications less on each example, which can reduce the impact that class imbalances can have.

After performing hyperparameter tuning, the results of the random forest classifier are lackluster when compared with that of Naive Bayes, but are significantly better than that of the support vector machine. The results can be found in Figure 4 on page 6.

	Precision	Recall	F1-Score	Support
Toxic	0.96	0.06	0.12	6090
Severely Toxic	0.00	0.00	0.00	367
Obscene	0.95	0.09	0.16	3691
Threat	0.00	0.00	0.00	211
Insult	0.67	0.01	0.03	3427
Identity Hate	0.00	0.00	0.00	712
Micro Average	0.93	0.05	0.10	14498
Macro Average	0.43	0.03	0.05	14498
Weighted Average	*0.80*	*0.05*	*0.10*	*14498*

Table 3: Tuned SVM on Test Set

	Precision	Recall	F1-Score	Support
Toxic	0.57	0.76	0.65	6090
Severely Toxic	0.23	0.08	0.12	367
Obscene	0.58	0.68	0.63	3691
Threat	0.33	0.05	0.09	211
Insult	0.56	0.52	0.54	3427
Identity Hate	0.57	0.12	0.20	712
Micro Average	0.57	0.62	0.59	14498
Macro Average	0.47	0.37	0.37	14498
Weighted Average	*0.56*	*0.62*	*0.57*	*14498*

Table 4: Tuned Random Forest on Test Set

5.5   Recurrent Neural Network (RNN) *Gabriel Britain*

The final model constructed was a recurrent neural network classifier. The use of the bag-of-words interpretation of comments loses a significant amount of information, like sentence structure and word meaning. Recurrent neural networks, specifically LSTMs, which are designed for long sequences, have demonstrated great promise in text classification. Another promising technology that has emerged through neural networks are word embeddings, which have been shown to model word semantic meaning in large corpi. We constructed a simple recurrent neural network consisting of two LSTM layers, a pooling layer, a dropout layer, and finally a densely-connected layer with a sigmoid activation function. The model was evaluated using a binary-crossentropy loss function after 10 epochs, and the results can be seen in Figure 5 on page 7. The architecture for the model can be seen in Figure 3 on page 8

	Precision	Recall	F1-Score	Support
Toxic	0.57	0.85	0.68	6090
Severely Toxic	0.34	0.48	0.40	367
Obscene	0.60	0.80	0.68	3691
Threat	0.00	0.00	0.00	211
Insult	0.52	0.72	0.61	3427
Identity Hate	0.67	0.22	0.34	712
Micro Average	0.56	0.75	0.64	14498
Macro Average	0.45	0.51	0.45	14498
Weighted Average	*0.56*	*0.75*	*0.63*	*14498*

Table 5: RNN on Test Set

6   Results Analysis *Gabriel Britain, Jonathan Innis*

After completing all of our tuned hyperparameter runs, we extracted features from the Naive Bayes classifier, and viewed the words and phrases that contributed most to each classification. Selected results can be seen in Figure 4 on page 9.

7   Conclusions *Gabriel Britain, Jonathan Innis*

After evaluating the results of each model, it can be concluded that the Naive Bayes classifier has the best performance; however, we believe that this is largely due to the significant imbalance of classes and the size of data. If more data was available, it makes sense that that the RNN would greatly benefit, and possibly surpass the Naive Bayes model. After this analysis, it can be concluded that the RNN with custom word embeddings would benefit the most from more data. Currently, the RNN built uses pre-trained word embeddings, but if provided more data, using custom word embeddings could provide better results.
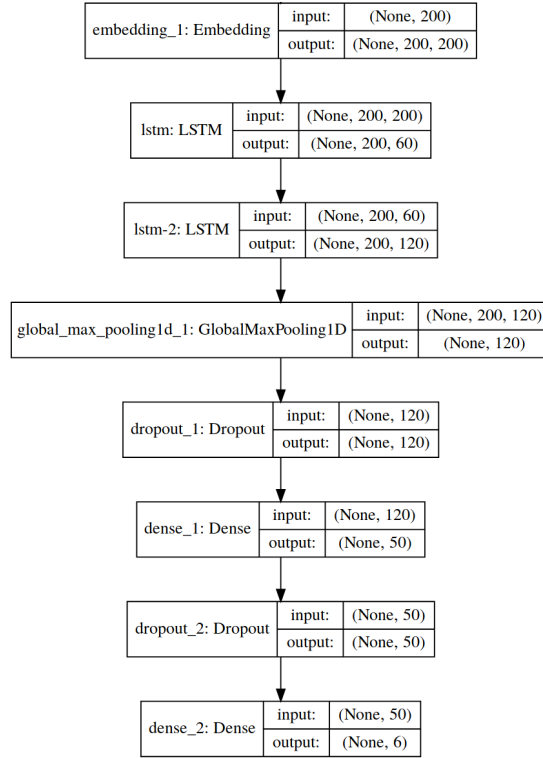
Figure 3: RNN Model Architecture

References

[1] Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, Greg S. Corrado, Andy Davis, Jeffrey Dean, Matthieu Devin, Sanjay Ghemawat, Ian Goodfellow, Andrew Harp, Geoffrey Irving, Michael Isard, Yangqing Jia, Rafal Jozefowicz, Lukasz Kaiser, Manjunath Kudlur, Josh Levenberg, Dandelion Mané, Rajat Monga, Sherry Moore, Derek Murray, Chris Olah, Mike Schuster, Jonathon Shlens, Benoit Steiner, Ilya Sutskever, Kunal Talwar, Paul Tucker, Vincent Vanhoucke, Vijay Vasudevan, Fernanda Viégas, Oriol Vinyals, Pete Warden, Martin Wattenberg, Martin Wicke, Yuan Yu, and Xiaoqiang Zheng. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.

[2] Lars Buitinck, Gilles Louppe, Mathieu Blondel, Fabian Pedregosa, Andreas Mueller, Olivier Grisel, Vlad Niculae, Peter Prettenhofer, Alexandre Gramfort, Jaques Grobler, Robert Layton, Jake VanderPlas, Arnaud Joly, Brian Holt, and Gaël Varoquaux. API design for machine learning software: expe-

toxic

```
yammer
follarte
fuckyourself
crackhead
```

severe toxic

```
stomes
cartuchos
ancestryearly
```

obscene

```
achmed
sexmist
britch
zigabo
```

threat

```
m45terbate
ma5terb8
masterbat3
zigabo
```

insult

```
crackhead
libtard
suberbia
```

identity hate

```
gomnna
zigabo
nebracka
```

Figure 4: Most Valuable Features

riences from the scikit-learn project. In *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, pages 108–122, 2013.

[3] François Chollet et al. Keras. https://keras.io, 2015.

[4] J. D. Hunter. Matplotlib: A 2d graphics environment. *Computing In Science & Engineering*, 9(3):90–95, 2007.

[5] Wes McKinney. Data structures for statistical computing in python. In Stéfan van der Walt and Jarrod Millman, editors, *Proceedings of the 9th Python in Science Conference*, pages 51 – 56, 2010.

[6] Jeffrey Pennington, Richard Socher, and Christopher D. Manning. Glove: Global vectors for word representation. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, 2014.

[7] Ellery Wulczyn, Nithum Thain, and Lucas Dixon. Ex machina: Personal attacks seen at scale. In *Proceedings of the 26th International Conference on World Wide Web*, WWW '17, pages 1391–1399, Republic and Canton of Geneva, Switzerland, 2017. International World Wide Web Conferences Steering Committee.